

Hidden layer produces outputs \mathbf{z}^H and \mathbf{u}^H of dimension N_h .

Output layer produces outputs \mathbf{z}^O and \mathbf{u}^O of dimension N_o .

Hidden layer: $z_j^H = \sum_{k=1}^{N_i} W_{jk}^H x_k + b_j^H, \quad u_j^H = g_{\text{act}}(z_j^H), \quad j = 1, \dots, N_h$

Output layer: $z_j^O = \sum_{k=1}^{N_h} W_{jk}^O u_k^H + b_j^O, \quad u^O = g_{\text{out}}(\mathbf{z}^O). \quad j = 1, \dots, N_o.$

Activation Functions:

– Hard threshold:

$$g_{\text{act}}(z) = \begin{cases} 1, & \text{if } z \geq 0 \\ 0, & \text{if } z < 0. \end{cases}$$

– Sigmoid: $g_{\text{act}}(z) = 1/(1 + e^{-z})$.

– Rectified linear unit (ReLU): $g(z) = \max\{0, z\}$.

Output Functions:

Binary classification: In this case, the response is $y = \{0, 1\}$. To predict the response, we typically take $N_o = 1$ and the output u^O is the class probability:

$$P(y = 1 | \mathbf{x}) = u^O = g_{\text{out}}(z^O) = \frac{1}{1 + e^{-z^O}}. \quad (3)$$

The variable z^O is called the *logit* and the output mapping (3) is a sigmoid. We can also use z^O to make a hard decision by selecting the most likely class:

$$\hat{y} = \begin{cases} 1 & z^O \geq 0 \\ 0 & z^O < 0. \end{cases} \quad (4)$$

K -class classification: In this case, the target is a class label $y = 1, \dots, K$. We typically take $N_o = K$ and use the soft-max function for the class probability:

$$P(y = k | \mathbf{x}) = u_k^O = g_{\text{out},k}(z^O) = \frac{e^{z_k^O}}{\sum_{\ell=1}^K e^{z_\ell^O}}. \quad (5)$$

Again, we can make a hard decision on the class label by selecting the highest logit score:

$$\hat{y} = \arg \max_{k=1, \dots, K} z_k^O. \quad (6)$$

Regression: In this case, one is trying to predict $\mathbf{y} \in \mathbb{R}^d$, where d is the number of variables to predict and each component y_j is typically continuous-valued. For this problem, we take $N_o = d$ and \mathbf{u}^O is the prediction of \mathbf{y} ,

$$\hat{\mathbf{y}} = \mathbf{u}^O = g_{\text{out}}(\mathbf{z}^O) = \mathbf{z}^O. \quad (7)$$

In (7), we will either say there is no activation or an *identity* activation.

When $N_o = 1$

Output layer: $z^O = \sum_{k=1}^{N_h} W_k^H u_k^H + b^O, \quad \hat{y} = g_{\text{out}}(z^O).$

$$\mathbf{z}^H = \mathbf{W}^H \mathbf{x} + \mathbf{b}^H, \quad \mathbf{u}^H = g_{\text{act}}(\mathbf{z}^H),$$

$$\mathbf{z}^O = \mathbf{W}^O \mathbf{u}^H + \mathbf{b}^O, \quad \mathbf{u}^O = g_{\text{out}}(\mathbf{z}^O).$$

Hidden layer: $z_{ij}^H = \sum_{k=1}^{N_i} W_{jk}^H x_{ik} + b_j^H, \quad u_{ij}^H = g_{\text{act}}(z_{ij}^H), \quad j = 1, \dots, N_h$

Output layer: $z_{ij}^O = \sum_{k=1}^{N_h} W_{jk}^O u_{ik}^H + b_j^O, \quad u_j^O = g_{\text{out}}(z_i^O), \quad j = 1, \dots, N_o.$

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_N^T \end{bmatrix} = \begin{bmatrix} x_{11} & \cdots & x_{1,N_i} \\ \vdots & \vdots & \vdots \\ x_{N1} & \cdots & x_{N,N_i} \end{bmatrix}, \quad \mathbf{Z}^H = \begin{bmatrix} (\mathbf{z}_1^H)^T \\ \vdots \\ (\mathbf{z}_N^H)^T \end{bmatrix} = \begin{bmatrix} z_{11}^H & \cdots & z_{1,N_h}^H \\ \vdots & \vdots & \vdots \\ z_{N1}^H & \cdots & z_{N,N_h}^H \end{bmatrix}$$

$$\mathbf{J}^H = \begin{bmatrix} (\mathbf{u}_1^H)^T \\ \vdots \\ (\mathbf{u}_N^H)^T \end{bmatrix} = \begin{bmatrix} u_{11}^H & \cdots & u_{1,N_h}^H \\ \vdots & \vdots & \vdots \\ u_{N1}^H & \cdots & u_{N,N_h}^H \end{bmatrix}, \quad \mathbf{U}^O = \begin{bmatrix} (\mathbf{u}_1^O)^T \\ \vdots \\ (\mathbf{u}_N^O)^T \end{bmatrix} = \begin{bmatrix} u_{11}^O & \cdots & u_{1,N_o}^O \\ \vdots & \vdots & \vdots \\ u_{N1}^O & \cdots & u_{N,N_o}^O \end{bmatrix}$$

$$\mathbf{Z}^H = \mathbf{X} [\mathbf{W}^H]^T + \mathbf{B}^H, \quad \mathbf{U}^H = g_{\text{act}}(\mathbf{Z}^H) \quad \mathbf{B}^H = \begin{bmatrix} (\mathbf{b}^H)^T \\ \vdots \\ (\mathbf{b}^H)^T \end{bmatrix} \quad \mathbf{B}^O = \begin{bmatrix} (\mathbf{b}^O)^T \\ \vdots \\ (\mathbf{b}^O)^T \end{bmatrix}$$

$$\mathbf{Z}^O = \mathbf{U}^H [\mathbf{W}^O]^T + \mathbf{B}^O \quad \mathbf{U}^O = g_{\text{out}}(\mathbf{Z}^O)$$

2. Consider the data set with scalar features x_i and binary class labels $y_i = \pm 1$.

x_i	0	1.3	2.1	2.8	4.2	5.7
y_i	-1	-1	-1	1	-1	1

Consider a linear classifier for this data of the form,

$$\hat{y} = \begin{cases} 1 & z > 0 \\ -1 & z < 0, \end{cases} \quad z = x - t,$$

where t is a threshold. For each threshold t , let $J(t)$ denote the sum hinge loss,

$$J(t) = \sum_i \epsilon_i, \quad \epsilon_i = \max(0, 1 - y_i z_i).$$

- Write a short python program to plot $J(t)$ vs. t for 100 values of t in the interval $t \in [0, 5]$.
- Based on the plot, what is one value of t that minimizes $J(t)$.
- For the value of t in part (b), find the corresponding slack variables ϵ_i .
- Which samples i violate the margin ($\epsilon_i > 0$) and which samples i are misclassified ($\epsilon_i > 1$).

2. (a) You can compute and plot $J(t)$ with the following code. The figure is shown in Fig. 1.

```
x = np.array([0, 1.3, 2.1, 2.8, 4.2, 5.7])
y = np.array([-1, -1, -1, 1, -1, 1])
tvals = np.linspace(0, 6, 100)
J = []
for t in tvals:
    z = x - t
    Jt = np.sum(np.maximum(0, 1 - y*z))
    J.append(Jt)
J = np.array(J)

plt.plot(tvals, J)
plt.xlabel('t')
plt.ylabel('J(t)')
plt.grid()

plt.savefig('Jt.png')
```

- From Fig. 1, we see that $t = 4$ is a minimizer.
- For the value of t in part (b), find the corresponding slack variables ϵ_i . We can compute the slack variables by the python code:

```
t = 4
z = x - t
eps = np.maximum(0, 1 - y*z)
eps

>> array([ 0. ,  0. ,  0. ,  2.2,  1.2,  0. ])
```

Exponential Functions

$$\frac{d}{dx} (e^x) = e^x$$

$$\frac{d}{dx} (a^x) = a^x \ln a$$

$$\frac{d}{dx} (e^{g(x)}) = e^{g(x)} g'(x)$$

$$\frac{d}{dx} (a^{g(x)}) = \ln(a) a^{g(x)} g'(x)$$

Logarithmic Functions

$$\frac{d}{dx} (\ln x) = \frac{1}{x}, \quad x > 0$$

$$\frac{d}{dx} \ln(g(x)) = \frac{g'(x)}{g(x)}$$

$$\frac{d}{dx} (\log_a x) = \frac{1}{x \ln a}, \quad x > 0$$

$$\frac{d}{dx} (\log_a g(x)) = \frac{g'(x)}{g(x) \ln a}$$

2. Suppose that we are given a mini-batch of training data (\mathbf{x}_i, y_i) , $i = 1, \dots, N$, and we try to fit a model to the data of the form,

$$\mathbf{z}_i = \mathbf{W}\mathbf{x}_i + \mathbf{b}, \quad \hat{y}_i = \sum_{j=1}^M \alpha_j \exp(z_{ij}),$$

for unknown parameters $\theta = (\mathbf{W}, \mathbf{b}, \alpha)$. Here M is the dimension of the hidden variables \mathbf{z}_i . We use the squared error loss function,

$$L = \sum_{i=1}^N L_i, \quad L_i = (y_i - \hat{y}_i)^2.$$

- Add an intermediate variable $\mathbf{u}_i = \exp(\mathbf{z}_i)$, by which we mean $u_{ij} = \exp(z_{ij})$. Draw the computation graph showing the mapping from the inputs \mathbf{x}_i and the parameters θ to the loss function, L . Also, write the equations for each step in the computation graph.
- Write the back-propagation equations to obtain the gradients with respect to the parameters θ .

2. (a) If we substitute $\mathbf{u}_i = \exp(\mathbf{z}_i)$ we can write the mapping from the set of \mathbf{x}_i 's to L as

$$\begin{aligned} \mathbf{z}_i &= \mathbf{W}\mathbf{x}_i + \mathbf{b} \\ \mathbf{u}_i &= \exp(\mathbf{z}_i), \quad \hat{y}_i = \alpha^\top \mathbf{u}_i \\ L &= \sum_{i=1}^N (y_i - \hat{y}_i)^2. \end{aligned}$$

The computation graph is shown in Fig. 2.

(b) We perform back-propagation in the following steps:

- $L \rightarrow \hat{y}$: The components of the gradient are:

$$\frac{\partial L}{\partial \hat{y}_i} = 2(\hat{y}_i - y_i).$$

- $\hat{y} \rightarrow \alpha, \mathbf{u}$: We have

$$\hat{y}_i = \alpha^\top \mathbf{u}_i = \sum_{j=1}^M \alpha_j u_{ij}.$$

Therefore,

$$\frac{\partial \hat{y}_i}{\partial u_{ij}} = \alpha_j, \quad \frac{\partial \hat{y}_i}{\partial \alpha_j} = u_{ij}.$$

Applying chain rule,

$$\begin{aligned} \frac{\partial L}{\partial \alpha_j} &= \sum_{i=1}^N \frac{\partial L}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial \alpha_j} = \sum_{i=1}^N \frac{\partial L}{\partial \hat{y}_i} u_{ij} \\ \frac{\partial L}{\partial u_{ij}} &= \frac{\partial L}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial u_{ij}} = \sum_{i=1}^N \frac{\partial L}{\partial \hat{y}_i} \alpha_j. \end{aligned}$$

- $\mathbf{u} \rightarrow \mathbf{z}$: Since $u_{ij} = \exp(z_{ij})$, we have the derivative,

$$\frac{\partial u_{ij}}{\partial z_{ij}} = \exp(z_{ij}).$$

Using chain rule:

$$\frac{\partial L}{\partial z_{ij}} = \frac{\partial L}{\partial u_{ij}} \frac{\partial u_{ij}}{\partial z_{ij}} = \frac{\partial L}{\partial u_{ij}} \exp(z_{ij}).$$

- $\mathbf{z} \rightarrow \mathbf{W}, \mathbf{b}$: Since $\mathbf{z}_i = \mathbf{W}\mathbf{x}_i + \mathbf{b}$, Therefore,

$$z_{ij} = \sum_k W_{jk} x_{ik} + b_j.$$

Hence, we have the derivatives,

$$\frac{\partial z_{ij}}{\partial W_{jk}} = u_{ik}, \quad \frac{\partial z_{ij}}{\partial b_j} = 1.$$

Applying chain rule,

$$\begin{aligned} \frac{\partial L}{\partial W_{jk}} &= \sum_i \frac{\partial L}{\partial z_{ij}} \frac{\partial z_{ij}}{\partial W_{jk}} = \sum_i \frac{\partial L}{\partial z_{ij}} u_{ik}, \\ \frac{\partial L}{\partial b_j} &= \sum_i \frac{\partial L}{\partial z_{ij}} \frac{\partial z_{ij}}{\partial b_j} = \sum_i \frac{\partial L}{\partial z_{ij}}. \end{aligned}$$

2. Consider the data set for four points with scalar features x_i and binary class labels $y_i = 0, 1$.

x_i	0	1	3	5
y_i	0	0	1	0

- Find a neural network with $N_h = 2$ units, $N_o = 1$ output units such that $\hat{y}_i = y_i$ on all four data points. Use a hard threshold activation function (2) and threshold output function (4). State the weights and biases used in each layer.
- Compute the values of \hat{y}_i and all the intermediate variables z_i^h , u_i^h and z_i^o for each sample $x = x_i$.
- Now suppose we are given a new sample, $x = 3.5$. What does the network predict as \hat{y} ?

2. (a) We can use a network similar in structure to the previous problem. In the hidden layer, we want to find features that can distinguish between $x = 3$ which belongs to class $y = 1$, and $x = \{0, 1, 5\}$ which belong to class $y = 0$. There are lots of choices, but we will extract two features: whether $x \geq 2$ and $x \geq 4$. We can do this with the linear transforms:

$$\mathbf{W}^h = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}, \quad \mathbf{b}^h = \begin{bmatrix} -2 \\ -4 \end{bmatrix},$$

	Training				Test
x_i	0	1	3	5	3.5
$z_{i1}^h = x - 2$	-2	-1	1	3	1.5
$z_{i2}^h = x - 2$	-4	-3	-1	1	0.5
$u_{i1}^h = g_{\text{act}}(z_{i1}^h)$	0	0	1	1	1
$u_{i2}^h = g_{\text{act}}(z_{i2}^h)$	0	0	0	1	0
$z_i^o = u_{i1}^h - 2u_{i2}^h - 0.5$	-0.5	-0.5	0.5	-1.5	0.5
$\hat{y}_i = g_{\text{out}}(z_i^o)$	0	0	1	0	1
y_i	0	0	1	0	

Table 1: Computations for the neural network in Problem 2

so that

$$\mathbf{z}^h = \mathbf{W}^h \mathbf{x} + \mathbf{b}^h = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} x + \begin{bmatrix} -2 \\ -4 \end{bmatrix} = \begin{bmatrix} x - 2 \\ x - 4 \end{bmatrix}.$$

Then, the activation outputs in the hidden layer are

$$\mathbf{u}^h = g_{\text{act}}(\mathbf{z}^h) = \begin{bmatrix} g_{\text{act}}(x - 2) \\ g_{\text{act}}(x - 4) \end{bmatrix} = \begin{bmatrix} \mathbb{1}_{\{x \geq 2\}} \\ \mathbb{1}_{\{x \geq 4\}} \end{bmatrix}.$$

Now we need to combine these functions in a way that they are negative for samples x_i when $y_i = 0$ and positive on sample x_i when $y_i = 1$. Similar to the previous problem, take

$$\mathbf{W}^o = [1, -2], \quad \mathbf{b}^o = -0.5.$$

Then,

$$\begin{aligned} z^o &= \mathbf{W}^o \mathbf{u}^h + \mathbf{b}^o = [1, -2] \begin{bmatrix} \mathbb{1}_{\{x \geq 2\}} \\ \mathbb{1}_{\{x \geq 4\}} \end{bmatrix} - 0.5 \\ &= \mathbb{1}_{\{x \geq 2\}} - 2\mathbb{1}_{\{x \geq 4\}} - 0.5 = \begin{cases} -0.5 & \text{when } x < 2 \\ 0.5 & \text{when } x \in [2, 4) \\ -1.5 & \text{when } x \geq 4. \end{cases} \end{aligned}$$

The predicted output \hat{y} is

$$\hat{y} = \mathbb{1}_{\{z^o \geq 0\}} = \begin{cases} 1 & \text{if } x \in [2, 4) \\ 0 & \text{else} \end{cases}$$

This matches the training data.

- Table 1 computes the values for all intermediate variables for the four training data points. We can see that $\hat{y}_i = y_i$ for all four points.
- The value for $x = 3.5$ is shown in the final column of Table 1. For this value of $x = 3.5$, we get $\hat{y} = 1$.